

Operating Systems


Recitation 8

Plan

- Thread synchronization
 - Lock & event in Linux
 - Mutex & Condition variable
 - What happens “under the hood”

Reminder

- Threads get scheduled by scheduler in kernel
- Preemptive multitasking:
 - OS decides when a thread will get its CPU time slot
- Context-switch without warning



Things can get ugly

```
void* mythread(void *arg) {  
    char *letter = arg;  
    int i;  
    printf("%s: begin\n", letter);  
    for (i = 0; i < TEN_MILLION; i++) {  
        balance = balance + 1;  
    }  
    printf("%s: done\n", letter);  
    return NULL;  
}
```

Code example

Things can get ugly

- **balance = balance + 1;**

- What really happens here? C translated to assembly
- In Intel x86* processor:

```
mov EAX, 0x8049a1c    // copy balance value
add EAX, 0x1          // increment by one
mov 0x8049a1c, EAX    // copy back
```

* <http://www.cs.virginia.edu/~evans/cs216/guides/x86.html>

Race condition

thread 1	thread 2	balance
mov EAX, 0x8049a1c		0
	mov EAX, 0x8049a1c	0
	add EAX, 0x1	0
	mov 0x8049a1c, EAX	1
add EAX, 0x1		1
mov 0x8049a1c, EAX		1 (not 2!)

- Result depends on the timing execution of the code
- Can get different result every time!

Let's take a closer look


- **balance = balance + 1;**

```
mov EAX, 0x8049a1c    // copy balance value
add EAX, 0x1           // increment by one
mov 0x8049a1c, EAX    // copy back
```

- What we have here is a **critical section**
- We need to access it with “atomicity”, force it to be executed **as a unit**
 - guaranteed to finish without other threads running the same code (at same time)

Synchronization for dummies

- Crossing our fingers ☹️
- Thread will run only after other thread ended ☹️
- **Delays** ☹️



Kernel sync mechanisms

- Atomic operations at **CPU level!**
 - Supported in instruction set
 - Example: intel x86 supports *atom_inc*, *atom_dec*, *atom_add*, *atom_sub*
- Interrupt disabling during critical section
 - No clock interrupt...
 - Hurts performance badly (unfair)
- Spinlocks (busy-wait)
 - Much more common
 - Saves context-switch
 - Useful only for very short periods

Linux Synchronization

- Main mechanism implemented in Linux for user-level synchronization
 - **Mutex** (lock)
 - **Condition variable** (signal)
- Windows
 - Critical section/Mutex
 - CreateEvent/WaitForSingleObject

Mutex

- Mutual exclusion
- Only **one** thread can hold it **lock()**ed
 - Others trying to `lock()` block until owner decides to free

- *Example:*

```
lock()
```

```
    //do some critical section code
```

```
    printf("hello")
```

```
unlock()
```

Mutex formal requirements

- Mutual exclusion
 - Only **one** thread can hold it lock()ed – can't have two in same critical section
- Progress (deadlock freedom)
 - **Some** thread eventually enters critical section
- Starvation freedom
 - Thread won't starve, and will **eventually** enter critical section

It's all in your head!

- Always remember when programming with mutexes:
it's a logical concept
- The protection of variables and code sections exists only *in your head*
- If you don't **consistently** protect shared variables/critical code with a mutex, bad things will happen
- OS only provides the mechanism - you are the user!



Mutex API

```
#include <pthread.h>
```

- Creation:

```
int pthread_mutex_init(  
    pthread_mutex_t *mutex,  
    const pthread_mutex_attr_t *mutexattr);
```

- Destruction:

```
int pthread_mutex_destroy(  
    pthread_mutex_t *mutex);
```

Mutex API - Example

```
pthread_mutex_t lock;

int main() {
    if (pthread_mutex_init(&lock, NULL) != 0) {
        printf("mutex init failed\n");
        return 1;
    }
    /* ... code here ... */
    pthread_mutex_destroy(&lock);
}
```

Locking

```
int pthread_mutex_lock(  
    pthread_mutex_t *mutex);
```

- If mutex is not free, block until it frees

```
int pthread_mutex_trylock(  
    pthread_mutex_t *mutex);
```

- If mutex is not free, fail

```
int pthread_mutex_unlock(  
    pthread_mutex_t *mutex);
```

- Free locked mutex

Under the hood

- Shared global variable acts as a 'lock'
- Initially 'unlocked'
 - `int mutex = 0;`
- Before entering critical section, a task 'locks' the mutex
 - `mutex = 1;`
- When done with critical section, 'unlocks' the mutex
 - `mutex = 0;`
- While mutex is "locked", no other task can enter critical section
- **What's the problem?**

Under the hood

- Special mutex variable needs to be accessed atomically
- Reasonable solution - hardware support
- One example (from the past):

testandset <address>, rnew, rold

- Special atomic operation

```
int TestAndSet(int *lock, int new) {  
    int old = *lock; // save old value of &lock in memory  
    *lock = new;      // set new value  
    return old;       // return old value  
}
```

Simple implementation

```
void init() {  
    // 0 means lock is available, 1 means held by a thread  
    flag = 0;  
}  
void lock() {  
    // busy-wait (do nothing)  
    // exits loop only when old value is 0 == not locked!  
    while (TestAndSet(&flag, 1) == 1) ;  
}  
void unlock() {  
    flag = 0;  
}
```

Simple implementation

```
void init() {  
    // 0 means lock is available, 1 means held by a thread  
    flag = 0;  
}  
void lock() {  
    // busy-wait (do nothing)  
    // exits loop only when old value is 0 == not locked!  
    while (TestAndSet(&flag, 1) == 1) ;  
}  
void unlock() {  
    flag = 0;  
}
```

That's a LOT of *spinning*!
Too many time-slices
wasted by scheduler on
threads in hopeless loop

Also possibly *starvation*!
Doesn't ensure all threads
will eventually acquire lock!

Less naive implementation

- Add `yield()` instruction

```
void init() {  
    flag = 0;  
}  
void lock() {  
    while (TestAndSet(&flag, 1) == 1)  
        yield(); // give up CPU on lock failure  
}  
void unlock() {  
    flag = 0;  
}
```

Not good enough

- Say we have 100 threads -
 - First thread locks, and gets preempted
 - 99 threads now try to `lock()`, fail and `yield()`
 - Still a LOT of context switching...
- And starvation...

More realistic implementation

- Implemented as struct with queue
 - Add thread to queue when lock unavailable
 - in `unlock()`, wake up one thread in queue
- A bit over-simplified
 - Also, mostly replaced by **Compare-and-Swap** (or other instructions)

Events

- Allow thread1 to inform thread2 on some event
 - Thread2 can sleep meanwhile
- Allow sync. access to sensitive shared resource
- Extension to mutex

Example: simple queue

- thread1 enqueues, thread2 dequeues
- Without sync. access:
 - Both threads may change data together
 - Thread1 insertion not safe (memory addresses...)
 - Thread2 won't know when to deq (memory addresses, polling...)

Condition Variables (1)

- Allow thread to sleep-wait() on event

```
int pthread_cond_init(  
    pthread_cond_t *cond,  
    pthread_condattr_t *cond_attr);
```

```
int pthread_cond_destroy(pthread_cond_t *cond);
```

- Initialize/destroy condition variable object
 - cond_attr = NULL is default
- Destroy fails if threads are waiting

Condition Variables (2)

```
int pthread_cond_wait(  
    pthread_cond_t *cond,  
    pthread_mutex_t *mutex);
```

- Wait() on condition variable
- **Must have mutex already locked!**
- On success releases mutex and puts thread to sleep
- Several threads can wait()
 - But only one wakes up...

Condition Variables (3)

```
int pthread_cond_signal(pthread_cond_t *cond);
```

- Signal a single wait()ing thread to wake up
- Choice of awakened thread is arbitrary
- Notice – **no mutex**

Back to queue example

```
item dequeue() {  
    pthread_mutex_lock(&qlock);  
    while <queue is empty>  
        pthread_cond_wait(&notEmpty, &qlock);  
    /* ... remove item from queue ... */  
    pthread_mutex_unlock(&qlock);  
    /* .. return removed item */  
}
```

Why **while**?

Back to queue example

```
pthread_mutex_t qlock;  
pthread_cond_t notEmpty;  
/* ... initialization code ... */  
void enqueue(item x) {  
    pthread_mutex_lock(&qlock);  
    /* ... add x to queue ... */  
    pthread_cond_signal(&notEmpty);  
    pthread_mutex_unlock(&qlock);  
}
```

Another example: producer/consumer

- Thread 1 “produces” elements
 - Element counter
 - “Element to consume” variable
- Consumer threads “consume” elements
 - Wait on “Element to consume” variable
 - “consumes” it and notifies producer it’s ready for more

Code example